

VERSION 2.1 — TEMPLATE CATALOGUE

PageCraft — Daily Engineering Timeline · amendment issued D2 · R5 / AI (Hanish)

Change bars on page(s): 3, 4 — Pragna · Discovery + Templates column, and two AI-column cells

Version 2.1 supersedes 2.0. Version 2.0's generation-architecture amendment and version 1.1's billing gate are unchanged and apply in full.

What this amendment does

No day moves. No milestone moves. No engineer gains or loses days. Eight cells change in Pragna's column, two are added in Hanish's, and one is amended for automation.

The catalogue grows from **25 hand-authored templates** to **≥80 curated compositions plus unlimited generation**. A template becomes a saved composition rather than a folder of files.

25 remains the D20 launch commitment

6–10 curated by D5 · 20–25 by D10 — **matching the existing commitment exactly** · 45–60 by D15 · 60–90 by D20 as a stretch. The MVP is not placed at risk under any outcome, because everything above 25 is produced by a pipeline rather than by hand.

Why this is not a scope increase

Hand-authored, 90 templates is roughly 45 engineer-days and cannot be delivered. As compositions, a template is a short JSON document: the batch builder generates 90 drafts, and a human curates each in about fifteen minutes.

	25 hand-built	90 curated
Design work	12.5 days	~3 days
New tooling	—	~3 days
Total	~12.5 d	~6 d

Pragna's effort is approximately neutral, shifted from making pages to making the gallery good — search, filters, ranking, empty states. A strong gallery over 90 templates beats a weak gallery over 25.

Dependency

Curation cannot begin until the section library and the batch builder exist — Stage 2 of the v2.0 staging, around D6–D10. **The curated catalogue therefore lands in week 3, not week 1.** The D5 and D10 template commitments are met by curating the first drafts and are unchanged in number.

Contract note

Correction to version 2.0

Version 2.0 stated that the **template schema contract is unaffected**. Under this amendment that is no longer true — template storage changes from a file set to a composition. All four Day-1 frozen contracts are now in scope.

Hard constraint

Curated templates may use images **only** from the asset import pipeline. Ninety templates at roughly four images each is around 360 images; a manual licence audit at that scale is not credible, and attribution is specified as a system property rather than a user responsibility.

Capacity

Catalogue construction is approximately **900 provider requests, one-off** — impossible on the free tier, trivial on the paid tier. A build cost, not a per-user cost. It reinforces rather than alters the D6 billing gate. Serving a curated template to a user costs zero requests, as today.

Corrections to the grid

Eight cells in Pragna · Discovery + Templates. Two added in Hanish · AI. One amended for automation.

Day	Column	SUPERSEDED (as printed in the body)	REPLACE WITH
D4	Pragna	Template schema in code; source + normalise first templates	Template schema in code — a template is a composition. Source and normalise the first curated drafts.
D5	Pragna	Reach 10 templates w/ licences; intent → gallery clickable	Reach 6–10 curated compositions w/ licences (attribution automatic via the asset pipeline); intent → gallery clickable. <i>Milestone commitment unchanged.</i>
D7	Pragna	Filter chips + URL state	Filter chips + URL state. Add vertical search with alias matching — "dentist" matches <code>dental-clinic</code> .
D9	Pragna	Templates 11–15	Gallery search + category grouping (eleven groups). The catalogue itself is now batch-generated; the work here is discovery, not authoring.
D10	Pragna	Templates 16–18; filter 18 real templates	Curate to 20–25 ; filter over the real catalogue. Ranking weights with exact-vertical-match boost (+100 exact vertical, +30 group, +10 tone/palette, +1 per shared section). <i>D10 milestone commitment unchanged.</i>
D12	Hanish · AI	Prompt tuning	Prompt tuning; vertical profiles. Add: batch catalogue builder + thumbnail renderer (~1 day). <i>May slip to D16–D17 without affecting the 25-template launch commitment.</i>
D13	Pragna	Form-endpoint field	Form-endpoint field. Curation UI — internal only, rough is acceptable: review a draft, swap a variant, adjust copy, approve.
D15	Pragna	Mobile pass; Meera's path end to end	Mobile pass (including the gallery filter UI collapse); Meera's path end to end. Catalogue at 45–60 curated.
D16	Pragna	Templates 19–22	Curation pass toward 60. Gallery empty state and pagination / lazy thumbnail loading.
D17	Pragna	Templates 23–25	Curation pass toward 90 (stretch). Catalogue diversity check — no theme on more than ~15% of curated templates, no motion on more than ~25%; build fails otherwise.
D18	Pragna	Licence + attribution audit; thumbnail regen	Verify the automatic attribution pipeline across the catalogue; thumbnail regeneration is scripted and deterministic , not manual. Spot-check rather than audit.
D19	Pragna	Visual polish; onboarding empty states; copy review; Day-19 freeze	

Visual polish; onboarding empty states; copy review; Day-19 freeze. **Confirm the describe-your-own path is reachable without scrolling to the end of the gallery.**

Milestone exit conditions

D5 — "10 real templates" becomes "6–10 curated compositions". Number unchanged in spirit and in commitment. **D10** — "filter 18 real templates" becomes "filter the real catalogue, 20–25 curated". **D20** — add: catalogue coverage, the proportion of beta sign-ups whose business type is served by a curated template. Target $\geq 80\%$.

Not changed, and why

Every day, owner and milestone date. · **Preethi's, Adhyay's and Adithya's columns in full.** · **The 25-template launch commitment** — it is met by D10 exactly as originally planned; everything above is upside from the pipeline. · **The v1.1 billing gate.** The one-off ~900-request catalogue build reinforces it; it does not move it. · **Path B, the generate path.** It already exists in v2.0 and needs no new capability — only visibility inside the gallery.

PageCraft — Confidential, internal. Version 2.1 supersedes 2.0. · Full assessment: *Template Catalogue 90 — change assessment.*

VERSION 2.0 — GENERATION ARCHITECTURE

PageCraft — Daily Engineering Timeline · amendment issued D2 · R5 / AI (Hanish)

Change bars on page(s): 3, 4 (the D1–D20 grid)

Version 2.0 supersedes version 1.1. Version 1.1's billing-gate amendment is unchanged and applies in full.

What this amendment does

No day moves. No milestone moves. No owner changes. No engineer gains or loses days — all five remain at 20. Nine cells change content, one banner row is added, and one gate is proposed.

Status: PROPOSAL. Stage 0 (the D2 cell) touches no frozen contract and is not covered by the Day-1 gate. Stages 1–4 require whole-team agreement and are contingent on the D5 go/no-go.

Why

The current design cannot produce an appropriate site for a business type nobody hand-authored a template for. Hospital, law firm, dentist, veterinary clinic, NGO, university, architecture firm, logistics, yoga studio and music school all resolve to a generic result.

The AI generation path changes from *fill one whole-page shell with HTML* to *emit a page composition over a library of section components*. **C-04 is preserved and strengthened** — the model emits data, code emits structure.

Add — banner row, below the product-framing row

INSERT:

Generation architecture (proposed, D2). From Stage 1 onward the AI emits a **page composition** — registered section types, registered layout variants, an art direction and typed content fields — rather than assembled HTML. Human-authored section components supply all markup. `projects.render_mode` keeps template-forked projects on the existing path; both modes coexist. Templates (Path A) are untouched and remain the generation fallback.

Staged: **Stage 0** (D2, seam only, no contract impact) · **Stage 1** (D3–D5, in memory, no schema) · **Stage 2** (D6–D10, persisted) · **Stage 3** (D11–D15, profiles and patches) · **Stage 4** (D16–D18, harden). Stages 1–4 contingent on the D5 go/no-go.

Proposed new gate

End D3 · E2 + E4

Section registry and the first section components exist and render. Hanish's D4 spike and D5 milestone both depend on them.

This is the only new hard dependency the architecture introduces, and it lands inside week 1. Previously the AI column fed the Editor column and never the reverse. MVP requirement is six section types at one variant each — hours of work, not days. Adithya's unblocking reserve should be available to Preethi in week 1.

Corrections to the grid

Nine cells. No day, owner or milestone moves.

Day	Column	SUPERSEDED (as printed in the body)	REPLACE WITH
D2	Hanish · AI	Provider interface	Provider methods real — structured output. fillSection() returns typed content fields, not HTML (Stage 0 seam). Motion tokens and the entrance observer added to the shell.
D3	Preethi · Frontend / Editor	CodeMirror integration; open/edit files from VFS	CodeMirror integration; open/edit files from VFS. Section registry skeleton and the first two section components (hero, services — one variant each).
D4	Hanish · AI	Two-stage generation quality spike (10 prompts)	Two-stage generation quality spike (10 prompts) — spike assembles a composition and renders it through the section components. Corpus spans 10 verticals, including at least three with no hand-authored template.
D4	Preethi · Frontend / Editor	iframe preview wired to VFS; live re-render on edit	iframe preview wired to VFS; live re-render on edit. Composition renderer rendering to React . Four more section components.
D6	Hanish · AI	Generation prod-shape — streaming + Zod validation	Generation prod-shape — streaming + Zod validation over the composition. <code>page_versions</code> with <code>schemaVersion</code> . <code>render_mode = 'composition'</code> for newly generated projects.
D8	Hanish · AI	Scoped-edit endpoint	Scoped-edit endpoint returning a content patch scoped to one section instance, with a plain-language explanation.
D8	Preethi · Frontend / Editor	Diff-view component (render proposed change, accept/reject)	Change-summary component — renders the proposed change in plain language, accept/reject. No diff, hunk or commit vocabulary in user-facing copy (consumer-tool framing).
D9	Preethi · Frontend / Editor	Preview + iframe sandbox hardening	Preview + iframe sandbox hardening. Deterministic edit routes wired — reorder, show/hide, lock, variant change and restyle, all at zero provider cost.
D11	Preethi · Frontend / Editor	Content panel complete (repeatable lists)	Content panel generated from registered content schemas , including repeatable lists. Adding a section type requires no editor change. <i>Expected to be less work than the superseded cell, not more — the per-template panel it replaces does not need doing.</i>
D11	Hanish · AI	Quality pass on 30 prompts (part 1) [requires billing enabled at D5]	Quality pass on 30 verticals (part 1) [requires billing enabled at D5]. Adds the visual-diversity metric — distinct (theme, motion, variant-set) triples across the sample.
D12	Hanish · AI	Prompt tuning [requires billing enabled at D5]	

			Prompt tuning — plan-stage variant selection and fill-stage copy. Vertical-profile generation and caching. [requires billing enabled at D5]
D14	Hanish · AI	Deterministic ranking	Deterministic ranking (unchanged — still drives template fallback and the gallery). Art direction dials fully wired; themes 3 → 8.
D16	Hanish · AI	Final prompt tuning	Final prompt tuning. Motion budget and diversity checks in the composition validator.
D18	Preethi · Frontend / Editor	Keyboard navigation	Keyboard navigation. LCP and CLS audit; reduced-motion verification across the four-browser matrix.

Milestone exit conditions — wording only, dates and owners unchanged

D5 — add to the exit condition: evidence shall include at least three verticals for which no template exists, demonstrating that the generic-template outcome has been eliminated. The D5 gate also carries a second, separable decision: adopt Stages 1–4, or continue on the shell-based pipeline.

D15 — 90% of **30 verticals** produce a sensible, non-blank page; injection contained; visual-diversity metric recorded.

Not changed, and why

Every day, owner and milestone date. This amendment changes nine cells, adds one banner row and proposes one gate; it moves nothing. · **Pragna's entire column.** Path A — landing, intent capture, gallery, 25 templates, publish UI — is untouched, and templates remain the generation fallback. · **The Backend + Publish column and the Platform + Database column in full.** · **The v1.1 billing gate.** Generation remains 8–10 requests; the D11, D12, D15 and D18 evaluation runs still require the paid tier from D6. Owner: Adithya.

Capacity note

Under this architecture the dominant *editing* operations — reorder, show/hide, add, remove, change variant, restyle and direct content edits — consume **zero** provider requests and are excluded from quota accounting. This does not remove the need for billing before D6, which is driven by the evaluation runs, but it materially improves the position of real beta users against the shared daily ceiling from D20.

VERSION 1.1 - MISSING BILLING GATE ADDED

PageCraft — Daily Engineering Timeline · amendment issued D5 · R5 / AI (Hanish)
Change bars on page(s): 1, 2

Version 1.1 supersedes version 1.0. This amendment adds one task and flags four that depend on it.

What the D4 spike changed, and what this grid never absorbed

The D4 generation spike measured Gemini’s free tier at **20 requests/day** and **5/minute**, per project per model — about **2 full generations per day** shared across all five engineers and every beta user. The capacity errata issued at D5 concluded that **Gemini billing must be enabled before D6**, and named Adithya as owner.

That conclusion was applied to the AI schedule. It was never applied here. **This grid contains no task, on any day, in any column, for enabling billing.** Week 0 lists "billing alerts" — an alerting task, not a provisioning one.

The consequence is not a missing line item. Four AI-column tasks in this grid cannot run at all on the free tier, because each needs more requests in a day than the tier allows:

Day	Task as written	Requests needed	Free tier allows
D11	Quality pass on 30 prompts (part 1)	~210	20/day — cannot run
D12	Prompt tuning (regression re-run)	~210 per pass	cannot run
D15	90% of 30 prompts sensible	~210	cannot run
D18	Verify caps under real load	load test	cannot run

The D15 milestone — *"End to end live"* — inherits the same dependency, and the PRD's 90%-of-30-prompts success metric cannot be evidenced without it.

Corrections to the grid

Two cell changes and four annotations. No day, owner or milestone moves.

1 · Adithya · Platform + Database — D5 cell (add)

SUPERSEDED (as printed in the body):

Sentry + PostHog; verify sign-in by email & GitHub onto one account

REPLACE WITH:

Sentry + PostHog; verify sign-in by email & GitHub onto one account; **ENABLE GEMINI BILLING — hard gate, blocks D6 onward**. Est. a few dollars/month at beta scale (2.5 Flash: ~\$0.30/M input, ~\$2.50/M output). Enabling it also makes the PRD's 'never used for training' promise true from D6 rather than D20.

2 · Week 0 · Pre-launch prep (clarify)

SUPERSEDED (as printed in the body):

... stand up accounts (GitHub org+test org, Vercel, Supabase, Upstash, Sentry, PostHog, Unsplash, LLM) · billing alerts ...

REPLACE WITH:

... stand up accounts (...) · billing alerts **and a payment method on the Gemini project, so billing can be switched on at D5 without a procurement delay** ...

3 · Hanish · AI — D11, D12, D15, D18 (annotate)

SUPERSEDED (as printed in the body):

Quality pass on 30 prompts (part 1) · Prompt tuning · 90% of 30 prompts sensible · Verify caps under real load

REPLACE WITH:

Each unchanged in scope, each marked **[requires billing enabled at D5]**. Contingency if billing is refused: the eval corpus shrinks to what 20 requests/day supports (~2 prompts), and the PRD's 90%-of-30-prompts success metric must be renegotiated or dropped. That is a product decision, not an engineering one — escalate at D5, not D11.

Not changed, and why

Every day, owner and milestone. This amendment adds one task and four annotations; it moves nothing. • **The D5 go/no-go.** It runs on 10 prompts, which the free tier supports across two days. It is D11 onward that breaks. • **The consumer-tool framing and product decisions.** Untouched.

PageCraft — Daily Engineering Timeline

MVP · 4 weeks (20 working days) · 5 engineers · Day-by-day decomposition of PRD \$10 weekly deliverables. Estimates for Day-1 re-sizing, not commitments.

Week 0 · Pre-launch prep (all/lead, ~0.5 day each): close open decisions · read terms (GitHub API, LLM provider, Vercel Hobby, Unsplash, DPDP) · stand up accounts (GitHub org+test org, Vercel, Supabase, Upstash, Sentry, PostHog, Unsplash, LLM) · billing alerts · ToS & privacy policy · domain · name a beta-user owner. **Exit:** nothing blocks Day 1.

Product framing (decided): CONSUMER-TOOL. The guided content panel is the default, primary surface for every user; the code view is a secondary, opt-in "Advanced" toggle, never required. GitHub/repo/commit/diff language never appears in user-facing copy — publish/dashboard/history screens use plain language ("your site," "changes saved," "version history," "suggested change"). GitHub OAuth stays deferred to publish only.

	Preethi · Frontend / Editor	Pragna · Discovery + Templates	Pragna + Adhyay · Backend + Publish	Adithya · Platform + Database	Hanish · AI	Milestone / Exit
Week 1 · Foundations						
D1	Day-1 contract review (/api + file-map types); scaffold editor shell layout	Design tokens + shadcn setup; review template/ content schema contract	Buffer / contract alignment with Adithya + Preethi	Run Day-1 freeze (contracts + shared TS types); stand up repo/CI/Vercel/Supabase	Contract review; prompt-harness scaffolding	
D2	File tree component + in-memory VFS data structure	Landing page	Octokit spike — repo creation	Schema + RLS + migrations	Provider interface	
D3	CodeMirror integration; open/edit files from VFS	Intent capture (8–10 category cards); gallery grid on stub data	Octokit spike — push files	Email auth	Classification endpoint	
D4	iframe preview wired to VFS; live re-render on edit	Template schema in code; source + normalise first templates	Persistence + project-CRUD API scaffolding (frozen contract)	GitHub auth + account linking	Two-stage generation quality spike (10 prompts)	
D5	Wire hard-coded files end to end; polish shell for Skeleton demo	Reach 10 templates w/ licences; intent → gallery clickable	Octokit spike proven — real repo in a real account	Sentry + PostHog; verify sign-in by email & GitHub onto one account	Go / no-go on generation quality — documented with evidence	D5 — Skeleton walkable: sign in → pick category → 10 real templates; editor renders/ edits hard-coded files
Week 2 · Core Loop						
D6	Wire persistence to Pragna+Adhyay file-persistence API (load/save project)	Gallery on real data	File-persistence API (real)	Rate limiting on AI routes	Generation prod-shape — streaming + Zod validation	
D7	Autosave loop + explicit save points	Filter chips + URL state	Project CRUD (real)	Daily spend caps on AI routes	One-shot repair + fallback-to-template	
D8	Diff-view component (render proposed change, accept/reject)	Content-panel v1 + site-metadata fields	Fork-a-template flow (copy files, project, initial commit via isomorphic-git)	Kill switch	Scoped-edit endpoint	
D9	Preview + iframe sandbox hardening	Templates 11–15	Commit history	Integrator / unblocking reserve	Cost logging	
D10	Integration: edits survive refresh; diff-view demo	Templates 16–18; filter 18 real templates	Fork flow produces a genuine first commit	Harden — AI endpoints un-abusable; unblocking reserve	Prompt produces real site <45s; failure path proven	D10 — Core loop closed: both paths land in editor; edit, save, reload, commit history persists
Week 3 · Publish						
D11	AI chat box UI shell inside editor	Content panel complete (repeatable lists)	DeployProvider interface; repo creation + collision handling	Support Pragna+Adhyay on OAuth scope issues	Quality pass on 30 prompts (part 1)	
D12	Wire AI chat to Hanish scoped-edit endpoint; render result as diff	Unsplash image library	Single-commit push; GitHub Pages enablement	Harden error handling across all routes	Prompt tuning	
D13	Auto-commit before AI edits (via Pragna+Adhyay commit API)	Form-endpoint field	Deploy polling; republish (update, not duplicate)	Playwright happy-path E2E test	Injection-containment tests	
D14	Version list UI + restore action	Publish + connect UI	Contact-forms endpoint	E2E — walk signup → live URL	Deterministic ranking	
D15	End-to-end AI edit flow; reject leaves file untouched	Mobile pass; Meera's path end to end	Meta tags; publish end to end — repo + live URL	Unblocking reserve; every failure path returns a real message	90% of 30 prompts sensible; injection contained	D15 — End to end live: sign in → pick/describe → edit → publish → repo in own GitHub + live URL
Week 4 · Harden & Launch						

V1.1 CORRECTED - SEE BILLING GATE AMENDMENT

	Preethi · Frontend / Editor	Pragna · Discovery + Templates	Pragna + Adhyay · Backend + Publish	Adithya · Platform + Database	Hanish · AI	Milestone / Exit
D16	Editor bug burn-down (P0/P1)	Templates 19–22	Deploy edge case — private / org repos	Security review (tokens, RLS, sandbox, scope audit)	Final prompt tuning	
D17	Large-file editor performance	Templates 23–25	Deploy edge case — name collision / Pages-off	Load sanity	Cost dashboard	
D18	Keyboard navigation	Licence + attribution audit; thumbnail regen	Republish reliability	Prod config; kill-switch drill	Verify caps under real load	
D19	Empty / loading / error states everywhere; Day-19 freeze	Visual polish; onboarding empty states; copy review; Day-19 freeze	Route hardening; Day-19 freeze	Day-19 freeze; unblocking reserve	Document prompt library; Day-19 freeze	
D20	Final editor QA; beta launch watch	Accessibility baseline; launch landing page	Final publish verification; beta launch watch	Beta launch — watch Sentry	Cost-per-user is a known number; beta launch watch	D20 — Beta launch to 20–30 recruited users: watch Sentry + funnel